

bloom-7b1-fine-tune-final

model: Bloom-7b1 (size: 25G)

dataset: alpaca (22M)

Python version: 3.8.10

Command: finetune.py


System Hardware: CPU count 32; GPU count 1; GPU type NVIDIA
A100-SXM4-80GB



epoch: 8

training time: 61h

Yajing Zhang

Section 1





25

train/learning_rate

— bloom7b1-finetuned-02



train/epoch

— bloom7b1-finetuned-02

8

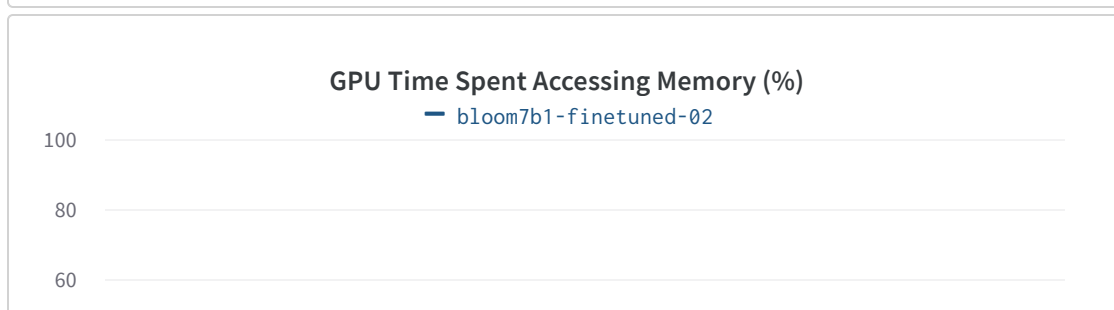






train/global_step

— bloom7b1-finetuned-02

3000







Created with ❤️ on Weights & Biases.

<https://wandb.ai/vmware/huggingface/reports/bloom-7b1-fine-tune-final--Vmlldzo0NDA4OTk2>